SSEE and the Growth Sector Re-examined: No Second Matter Component, and No S_8 Tension, on Raw Survey Data

Mike Edison Almeida Vallejo

Independent Researcher

mike.almeida1721@gmail.com

ORCID: 0009-0008-2195-7836

July 2026 (SSEE, Paper 6 of 10)

Abstract

We test the growth-of-structure sector of SSEE against *raw* survey data rather than against compressed statistics whose reduction assumes a ΛCDM background. The model has a single matter density, $\Omega_m = 0.308881$, entering the background and the growth of structure alike; the background is fixed by the algebra. Against the 357 raw $\xi_\pm$ points of KiDS-Legacy [Wright et al., 2025], with the amplitude free, the shear prefers $\log(10^{10}A_s) = 3.0255 \pm 0.0396$, 0.49σ from the 3.04483 that SSEE's own CMB fit supplies (under the identical pipeline with the Planck background, 1.01σ). Fixing the amplitude to the CMB value leaves the growth sector with *no free cosmological parameter* and costs nothing: $\chi^2_{\rm min} = 417.97$ against 418.34 with the amplitude free, $\Delta\mathrm{BIC} = 6.24$, and the nuisance parameters do not move to compensate. $S_8 = 0.8273$ is then an output of the algebra and the CMB, against a measured 0.8265 ± 0.0176. KiDS-Legacy predates this analysis by sixteen months and no temporal priority is claimed: the result is a consistency of A_s between epochs. We state against our own interest that the three χ^2 values compared span less than 1.0: cosmic shear does not discriminate between these backgrounds, and what separates them is the parameter count. On the earlier KiDS-1000 release (225 raw points, amplitude free), SSEE gives $\chi^2_{\rm min} = 265.44$ against 262.75 for ΛCDM run through the identical pipeline with its background free: four extra parameters buy $\Delta\chi^2 = 2.69$, less than the 4 they would return by chance ($\Delta\mathrm{BIC} = -19.0$ in favour of SSEE). Against the raw BOSS DR12 multipoles (222 points, one-loop LPT) the algebraic and Planck backgrounds fit equally well within noise, $\Delta\chi^2 = -2.03$ for SSEE at equal freedom; clustering does not constrain the amplitude. This paper also retracts a second matter sector and an associated particle ($m_\phi = 40.70$ eV) proposed in an earlier version: the sector was defined by subtracting $1 + w_0 \approx 0.160$

— a number of the dark-energy equation of state, not a density — from the matter density, and the S_8 tension it was built to close existed only in the compressed statistic, never in the raw data.

Contents

1 Introduction: why the growth sector needed re-examination **4**

2 The retraction: why the second sector is withdrawn **4**

2.1 The subtraction that never was a density 4

2.2 Falsifiability is not evidence of existence 5

2.3 Open problems closed by dissolution 5

3 Methodology: measuring against raw data, not compressed statistics **6**

3.1 The principle: combined data reports its own internal tensions 6

3.2 Test design: ingredients transported, not re-fitted 6

3.3 Raw KiDS-1000 cosmic shear likelihood 7

3.4 Raw BOSS DR12 redshift-space multipoles 7

4 S_8 against raw cosmic shear **8**

5 The same test against KiDS-Legacy **9**

5.1 Historical note: the particle's exclusion, for the record 12

6 $f\sigma_8$: growth rate against raw BOSS DR12 clustering **12**

6.1 Why the exploratory linear-Kaiser stage does not publish 14

7 Conclusions **15**

1 Introduction: why the growth sector needed re-examination

The background sector of SSEE — the algebraic derivation of H_0, ω_b, ω_c, n_s, w_0 and w_a from φ and π — has been tested against Planck 2018, DESI DR2 BAO, and Pantheon+ in Papers 1–4, 9–10. The growth-of-structure sector — how matter perturbations evolve once the background is fixed — is tested independently, against two late-time probes: weak gravitational lensing (cosmic shear, KiDS-1000) and redshift-space galaxy clustering (BOSS DR12). Both probes report their results as derived quantities: KiDS-1000 publishes $S_8 \equiv \sigma_8\sqrt{\Omega_m/0.3}$, and BOSS publishes $f\sigma_8(z)$, the growth rate times the amplitude of clustering.

Both derived quantities share a hidden dependency that matters for a model whose background departs from ΛCDM: converting a raw two-point statistic (shear correlation functions, or galaxy power-spectrum multipoles) into a compressed number requires assuming a fiducial cosmology — the Alcock–Paczynski effect [Alcock and Paczynski, 1979] for redshift-space clustering, and the non-linear matter power spectrum template for lensing. Both official pipelines assumed ΛCDM calibrated to Planck [Planck Collaboration, 2020]. Comparing a ΛCDM-conditioned number against a background that is *not* ΛCDM (SSEE has a Chevallier–Polarski–Linder [Chevallier and Polarski, 2001, Linder, 2003] equation of state $w_0 = -0.839950$, not -1) imports part of the ΛCDM background into the comparison itself.

This paper re-does both measurements on the raw data products — KiDS-1000 $\xi_\pm(\theta)$ (Asgari et al. 2021 [Asgari et al., 2021]) and BOSS DR12 power-spectrum multipoles $P_0, P_2, P_4(k)$ (Beutler et al. 2017 [Beutler et al., 2017]) — applying the model-dependent deformation (Alcock–Paczynski for BOSS; the non-linear matter power spectrum template for KiDS) to the *theory prediction*, never to the data. This keeps the measurement clean of any assumed background, at the cost of re-deriving both likelihoods from the covariance matrices up rather than reusing the compressed public numbers.

2 The retraction: why the second sector is withdrawn

2.1 The subtraction that never was a density

The two-sector construction, retired here, rested on

$$\Omega_{\phi\mathrm{DM}} \;=\; \Omega_{m,\mathrm{CMB}} - \Omega_{m,\mathrm{dyn}}, \qquad \Omega_{m,\mathrm{dyn}} \equiv 1 + w_0. \tag{1}$$

$\Omega_{m,\mathrm{CMB}} = \omega_m/h^2 = 0.308881$ is a measured density fraction, fixed by the CMB-calibrated background. $\Omega_{m,\mathrm{dyn}} = 1 + w_0 \approx 0.160050$ is, by construction, minus the dark-energy equation of state — an algebraic identity of the background sector's $w_0 = -T_R/M_V$, entering the Friedmann equations as an equation-of-state parameter, never as a density. Both quantities are dimensionless, so the subtraction is arithmetically well defined; but a

well-defined subtraction between a density and an equation-of-state number does not itself define a density, let alone a particle. The internal code comments that introduced this construction stated explicitly that $\Omega_{m,\mathrm{dyn}}$ "is NOT a matter density: it is the dynamical cold sector $(1 + w_0)$" in the same breath as summing it with $\Omega_{\phi\mathrm{DM}}$ to recover $\Omega_{m,\mathrm{CMB}}$ — an internal inconsistency once both statements are read together. $(1 + w_0$ is dimensionless and is, by construction, minus the dark-energy equation of state; the parenthetical gloss is ours, not the source comment's.)

2.2 Falsifiability is not evidence of existence

A canonical particle mass, $m_\phi = 40.70$ eV — retracted here — was derived from this subtraction via a fixed algebraic multiplier (retracted: $\mathrm{SOLAR}^2 \cdot \mathrm{KRYSTOS_V} = 594.28$) acting on $\Sigma m_\nu^{\mathrm{active}}$, explicitly to render the two-sector hypothesis falsifiable against free-streaming suppression in the matter power spectrum. Falsifiability is a necessary property of a scientific hypothesis; it is not evidence that the hypothesis's object exists. Once the subtraction that defined $\Omega_{\phi\mathrm{DM}}$ is shown not to correspond to a physical density, the particle built on top of it has nothing left to be made of, and is withdrawn together with it — independently of whatever bound the data would place on its mass. §4 reports that bound only as part of the historical record of how the retraction was reached, not as a surviving result.

2.3 Open problems closed by dissolution

Four open problems catalogued against the retired two-sector construction close by dissolution, following the precedent of OP-8 (the analogous Ω_m partition via the algebraic constant MIRA, retired 2026-07-09):

- **OP-9** (origin of the $\mathrm{SOLAR}^2 \cdot \mathrm{KRYSTOS}_V$ coefficient) — closes: there is no coefficient left to derive.
- **OP-10** (the "single free scalar field" claim, admitted as "technically violated" by the retired two-sector construction) — closes, and the single-field philosophy of SSEE is restored without qualification.
- **OP-11** (ξ, the non-minimal coupling of the second sector — an explicit additional free continuous parameter) — closes: the parameter is removed, not fitted.
- **OP-12** (T_ϕ, described in its own documentation as "bookkeeping, not a temperature derived from physics") — closes.

No new open problem is created by this retraction: what remains is $\Omega_m = 0.308881$, a single measured quantity, entering both the background and the growth-of-structure sector with no partition.

3 Methodology: measuring against raw data, not compressed statistics

3.1 The principle: combined data reports its own internal tensions

A dataset built by combining independent probes under a flexible model does not report fidelity to any one of them: it reports the point at which their mutual tensions are minimized, which is not the same thing. A model with enough freedom to fit each probe separately with a *different* set of nuisance values will inevitably report a lower χ^2 when its parameters are free to migrate between probes than when they are held fixed to what any single probe prefers —not because the combination is more accurate, but because it is averaging disagreement rather than resolving it. This is the generic mechanism behind the well-documented S_8 tension between Planck and cosmic shear surveys under ΛCDM [Di Valentino et al., 2021]: parameters fitted to the CMB [Planck Collaboration, 2020], carried over unchanged to lensing, land outside the region preferred by KiDS-1000 [Asgari et al., 2021, Heymans et al., 2021] and DES Y3 [Amon et al., 2022] alike.

SSEE is structurally unable to exploit this mechanism: its background $(\omega_b, \omega_c, h, n_s, w_0, w_a)$ is fixed by algebra, identical whether evaluated against the CMB, cosmic shear, or redshift-space clustering. The only cosmological parameter that could vary between probes is the primordial amplitude A_s (equivalently $\log A$). In SSEE it is counted once, at the CMB, and carried to every other probe fixed (§5).

3.2 Test design: ingredients transported, not re-fitted

To make this principle falsifiable rather than rhetorical, we perform the following test (a design due to M. Almeida): the three probes — Planck CMB, KiDS-1000 shear, BOSS DR12 clustering — are combined with the shared amplitude free, yielding the ingredient set the *combination* prefers. That same ingredient set is then evaluated against each probe *individually*, without re-fitting. Transporting parameters from a joint fit to an individual dataset always degrades χ^2 relative to a dedicated fit — that is arithmetic, true of any model, and proves nothing by itself. The diagnostic quantity is how much it degrades relative to what the probe's own statistical uncertainty predicts:

$$\Delta\chi^2(\text{probe}) \;=\; \chi^2\Big(\text{joint ingredients}\Big) - \chi^2\Big(\text{probe-fitted ingredients}\Big). \tag{2}$$

For SSEE, $\Delta\chi^2 \equiv 0$ by construction on every background parameter except A_s: there is no ingredient to transport, because none of $\omega_b, \omega_c, h, n_s, w_0, w_a$ are free to begin with. For ΛCDM, $\Delta\chi^2 > 0$ is measured, not assumed. Results in Table 1.

Table 1: Ingredient transport test (§3.2). $\Delta\chi^2$ evaluated per probe, joint-fit ingredients vs. probe-dedicated fit. ΛCDM: total loss 8.39 for 8 degrees of freedom (0.85σ), from the joint CMB+KiDS-Legacy+DESI fit (`lcdm_conjunta.json`).

	CMB (Planck)	shear (KiDS-Legacy)	BAO (DESI DR2)
ΛCDM, $\Delta\chi^2$	3.63	2.79	1.96
SSEE, $\Delta\chi^2$	0 (by construction)	0 (by construction)	0 (by construction)

3.3 Raw KiDS-1000 cosmic shear likelihood

We evaluate χ^2 against the 225 (of 270) scale-cut $\xi_\pm(\theta)$ data points from the KiDS-1000 Blind C release (Asgari et al. 2021 [Asgari et al., 2021]), using the flat-sky Limber approximation for the shear power spectrum C_ℓ^{ij} (GG+GI+II, NLA intrinsic alignment model, Bridle & King 2007 [Bridle and King, 2007]), HMcode-2015 non-linear matter power spectrum (Mead et al. 2015 [Mead et al., 2015]) with the one-parameter baryon feedback prescription used in the official KiDS-1000 analysis, and the official redshift-distribution SOM covariance and δ_c priors. The pipeline reproduces the published maximum-posterior likelihood to 0.4% (negative control: $\chi^2 = 261.27$ vs. 260.31 official, the latter from the public nested-sampling chain, MultiNest [Feroz et al., 2009]).

Eight nuisance parameters (baryon feedback amplitude, intrinsic alignment amplitude, five source-redshift shifts, additive c-term) are marginalized via MCMC (Cobaya [Torrado and Lewis, 2021], Metropolis sampler, 4 independent chains, Gelman–Rubin $R-1<0.03$), not minimized, for both SSEE (background fixed, $\log A$ free: 9 parameters total) and ΛCDM (background free: $\omega_b, \omega_c, h, n_s$; 13 parameters total, same nuisance treatment).

3.4 Raw BOSS DR12 redshift-space multipoles

We evaluate χ^2 against the full-shape power-spectrum multipoles $P_0, P_2, P_4(k)$ from BOSS DR12 (Beutler et al. 2017 [Beutler et al., 2017]), 3 redshift bins ($z = 0.38, 0.51, 0.61$) $\times$ 2 hemispheres (NGC/SGC), using the PATCHY mock covariance with Hartlap correction. The Alcock–Paczynski deformation is applied to the *model*, mapped onto the fiducial cosmology ($\Omega_m = 0.31$, $w = -1$, $H_0 = 100$, declared in the data headers) in which Beutler measured, so that the raw data vector is never touched. The survey window function is applied via the configuration-space convolution method of Wilson et al. (2017) [Wilson et al., 2017]: $P_\ell(k) \to \xi_\ell(s)$ (Hankel transform), mixed with the window multipoles W_0, W_2, W_4, W_6, W_8, transformed back. Redshift-space distortions are modelled at one loop in Lagrangian perturbation theory, in its effective-field-theory implementation (velocileptors; Baumann et al., 2012, Chen et al., 2020), to $k_{\rm max} = 0.20\,h\,{\rm Mpc}^{-1}$; an earlier linear-Kaiser pass is described in §6.1 and does not publish.

4 S_8 against raw cosmic shear

The SSEE chain (R3) samples nine free parameters — the primordial amplitude $\log A$ plus the eight KiDS-1000 nuisance parameters (the baryonic feedback amplitude $A_{\rm baryon}$, the intrinsic-alignment amplitude $A_{\rm IA}$, five photometric-redshift shifts δz_i, and the additive shear term δ_c) — against the 225 unmasked $\xi_\pm$ data points, with the background held fixed by the algebra ($\Omega_m = 0.308881$, ω_b, ω_c, h, n_s, w_0, w_a all determined, none of them fitted). Four independent MPI chains were run to a Gelman–Rubin convergence of $R - 1 = 0.019$, well inside the stopping criterion $R - 1 < 0.03$, giving 42711 effective samples after discarding the first 30% of each chain. The marginalized posterior is

$$\sigma_8 = 0.7449 \pm 0.0186, \qquad S_8 \equiv \sigma_8\sqrt{\Omega_m/0.3} = 0.7559 \pm 0.0189, \tag{3}$$

with a best-fit $\chi^2 = 265.44$ over 225 data points and 9 free parameters (216 degrees of freedom, $\chi^2/\mathrm{dof} = 1.229$). Both nuisance parameters that carry astrophysical meaning land in unremarkable places ($A_{\rm baryon} = 2.57 \pm 0.31$, $A_{\rm IA} = 0.54 \pm 0.29$), so the fit is not being rescued by a nuisance parameter running to the edge of its prior.

Where the model is actually tested: the fit to the raw data. S_8 is an output of a fit, not a measurement: quoting one model's S_8 against another's compares two *products* of two analyses. The quantity that confronts SSEE with something the model did not produce is the likelihood of the 225 raw $\xi_\pm$ points, so that is where we put the comparison (Table 2). Against the ΛCDM control run through the same pipeline, same mask, same covariance and same nuisance treatment,

$$\Delta\chi^2 \;\equiv\; \chi^2_{\rm SSEE} - \chi^2_{\Lambda\rm CDM} \;=\; 265.44 - 262.75 \;=\; 2.69, \qquad \Delta k = 4, \tag{4}$$

where Δk counts the parameters ΛCDM has and SSEE does not. Four additional free parameters that carried no information would still lower χ^2 by $\langle\Delta\chi^2\rangle = \Delta k = 4$ on average; here they lower it by 2.69, and $p(\Delta\chi^2 \geq 2.69 \,|\, \Delta k = 4) = 0.61$. The extra background freedom of ΛCDM is therefore not merely unnecessary on this dataset — it is statistically indistinguishable from absent. Penalising it explicitly gives $\Delta\mathrm{AIC} = -5.31$ and $\Delta\mathrm{BIC} = -19.0$, both favouring SSEE, the latter in the *strong* range on the Jeffreys scale.

The derived amplitudes are then a consequence rather than the test: $\Delta S_8 = 0.0012$, under a tenth of either error bar. Against the *published* KiDS-1000 value $S_8 = 0.759 \pm 0.024$ the offset is 0.10σ, but that published number is itself ΛCDM-derived, so this second comparison is informative about the pipeline rather than about the cosmology, and we do not rest any claim on it.

The one tension we do report, and it is shared. Both fits are formally poor in absolute terms: $\chi^2 = 265.44$ against an expectation of 216 ± 20.8 for SSEE and 262.75 against 212 ± 20.6 for ΛCDM, i.e. 2.4σ and 2.5σ high, with probabilities to exceed of 0.012

and 0.010. This is not a feature of either cosmology: the KiDS-1000 published best fit on these same points has $\chi^2 = 260.31$, which under any of the plausible effective-parameter counts gives a comparable $p \sim 0.01$–0.03. The excess is a property of the $\xi_\pm$ data vector — plausibly residual small-scale modelling or covariance calibration — and it is inherited identically by every model fitted to it. We flag it rather than absorb it, and note that our control's reproduction of the official value (262.75 vs. 260.31, 0.9%) is what demonstrates the excess is not introduced by our implementation.

What this does and does not claim. It would be wrong to read this as SSEE having carried a 3.5σ tension and then removing it. SSEE never carried that tension. At this level the model has exactly *one* free quantity, the primordial amplitude A_s; the background is fixed by the algebra and is identical in every fit reported here. What varies between determinations is only *which dataset is asked to fix* A_s, and the datasets do not agree with each other:

A_s determined by	$\log(10^{10}A_s)$	S_8	vs. published KiDS-1000[†]
raw KiDS-1000 $\xi_\pm$ (this work)	2.864 ± 0.050	0.7559 ± 0.0189	0.10σ
SSEE fit to raw Planck (full `plik`, B1 chains)	3.042 ± 0.013	0.8261 ± 0.0054	2.73σ
ΛCDM-fitted Planck value (borrowed)	3.045	0.8229 ± 0.0059	2.59σ

[†]The last column compares each row against the *published* KiDS-1000 value, which is a ΛCDM-derived number; it is given for orientation only and no claim here rests on it. The quantity that matters for the argument below is different: the internal disagreement between rows 1 and 2, which are the same model with the same background. In $\log(10^{10}A_s)$ that separation is $3.042 - 2.864 = 0.179 \pm 0.052$, i.e. $\mathbf{3.5}\sigma$.

The first two rows are the same model with the same background, differing only in the data used to set one number, and they sit 3.5σ apart. That separation is a disagreement *between the datasets*, not a property of SSEE: it persists when A_s is fixed by SSEE's *own* CMB fit rather than borrowed from ΛCDM, and it is the same Planck–shear discrepancy that ΛCDM also carries. **We do not claim to resolve it.**

What the earlier version of this paper did wrong was to attribute that data–data disagreement to the model — reporting it as "SSEE's S_8 tension" — and then to build a second matter sector and a particle to close it. The disagreement was never the model's to close, and the KiDS collaboration's own recalibration dissolved it in the data (§5).

5 The same test against KiDS-Legacy

The disagreement reported above was resolved by the KiDS collaboration itself, and not in the model's favour or against it: in the data. KiDS-Legacy [Wright et al., 2025] recalibrates the source redshift distributions and re-analyses the complete survey, and its Appendix I re-runs the KiDS-1000 analysis end to end with the new methods under blinding. The tension against Planck dissolves in steps, quoted there as signed consistencies in S_8 (negative meaning KiDS below Planck): $-2.39\sigma \to -1.42\sigma$ from the $n(z)$ calibration

alone, then -1.12σ, -0.85σ and -0.74σ as tomography, alignment model and sky area are updated. Their stated threshold is "the typical KiDS threshold for consistency of 2.3σ (equivalent to a one-sided probability to exceed of 0.01)"; every step with the new calibration falls below it.

Chronology, stated plainly. KiDS-Legacy appeared on arXiv on 2025 March 25, sixteen months before the analysis of §4 was written. **We claim no temporal priority over it.** What we report here is an internal consistency test of this model, run on the raw Legacy data vector, and the reader should weigh it as such.

We repeat the measurement of §4 on the 357 unmasked $\xi_\pm$ points of the Legacy release, in three configurations that differ only in what is held fixed. All three use the same likelihood, the same scale cuts (ξ_+: 2′–300′, ξ_-: 4′–300′) and the same eight nuisance parameters.

Configuration	background	A_s	free	$\chi^2_{\rm min}$	BIC
SSEE unified	algebraic	from CMB	8	417.97	464.99
SSEE	algebraic	free	9	418.34	471.24
ΛCDM-fixed	Planck 2018	free	9	418.95	471.85

Three results follow, and they answer different questions.

(i) The amplitude the shear asks for is the one the CMB fixes. With the background held by the algebra and $\log A$ free, the Legacy data prefer $\log(10^{10}A_s) = 3.0255 \pm 0.0396$, against the value SSEE's own CMB fit supplies, 3.04483: a separation of 0.49σ. Under the identical pipeline with the Planck background held fixed instead, the same data prefer 3.0069 ± 0.0368 against that background's own CMB value, 1.01σ away. The difference has a mechanism rather than being noise: SSEE's Ω_m is 2.00% lower than Planck's, and the shear asks it for an A_s 1.88% higher — less matter requires more amplitude to produce the same observed clustering. Each background *forces* a different A_s, and the one the algebraic background forces is the one its own CMB predicts. Both $\log A$ values come from the same 357 points and are therefore correlated; no clean significance can be quoted for their difference, and 1.01σ is not a tension either. What can be stated is the direction, and that it has an identified physical cause.

(ii) Fixing A_s costs nothing, and could have cost a great deal. Removing the only cosmological parameter from a fit removes its ability to accommodate: if either the algebraic background or the CMB amplitude were wrong, χ^2 rises with nowhere to hide it. It did not rise. With one parameter fewer, $\chi^2_{\rm min} = 417.97$ against 418.34, and $\Delta{\rm BIC} = 6.24$ in favour of the unified configuration (6.86 against ΛCDM-fixed). A second, independent indication points the same way: the nuisance parameters do not move to compensate. Had the fixed A_s been wrong, $\log T_{\rm AGN}$ and $A_{\rm scale}$ would have run to cover the gap; they shift by 0.066 and 0.10, well inside their own posterior widths (0.20 and 0.75). We note for completeness that $\Delta\chi^2$ is formally negative (-0.365), which cannot be a true improvement

— the fixed-parameter space is contained in the free one — and measures the MCMC sampling noise: eight dimensions are explored more efficiently than nine. The honest statement is that the cost is indistinguishable from zero, not that fixing improves the fit.

(iii) S_8 becomes a prediction. With both background and amplitude supplied, $S_8 = 0.8273$ is an output of the algebra and the CMB, with no free cosmological parameter behind it and therefore no posterior width of its own. The official Legacy $\xi_\pm$ posterior, computed from its own published chain as $\sigma_8\sqrt{\Omega_m/0.3}$, gives 0.8265 ± 0.0176. Two caveats bound what this establishes: the agreement in S_8 and the 0.49σ in $\log A$ are the same fact in two variables and do not compound, and the ± 0.0176 width means any prediction between 0.809 and 0.844 would have agreed. The test is not sharp. What it does establish is that the growth sector takes its amplitude from the CMB and describes cosmic shear without adjusting anything.

What this closes. The gap that the retracted light particle was introduced to fill is 0.49σ wide. There is no work left for it to do, and it is withdrawn as unnecessary — not excluded by a bound. Two things follow that are worth stating against our own interest. First, a particle that is not needed may still exist; this is not a constraint on its mass. Second, the three χ^2 values above span less than 1.0, so cosmic shear does not discriminate between these backgrounds: what separates them is the parameter count (BIC) and the cross-epoch consistency of A_s, neither of which is a goodness-of-fit statement. **The result is not that SSEE describes cosmic shear better; it is that SSEE describes it equally well without being allowed to fit it.**

Table 2: Fit to the *raw* KiDS-1000 $\xi_\pm$ data vector (225 points), primordial amplitude free in both models, MCMC-marginalized. The test is the left block: the likelihood of data neither model produced. The S_8 block is derived output, shown for continuity with the literature. p is the probability to exceed of $\chi^2_{\rm min}$ at the stated dof.

	k	dof	$\chi^2_{\rm min}$	χ^2/dof	p	σ_8	S_8
ΛCDM (background free)	13	212	262.75	1.239	0.010	0.875 ± 0.089	0.7571 ± 0.0194
SSEE (background fixed)	9	216	265.44	1.229	0.012	0.7449 ± 0.0186	0.7559 ± 0.0189
$\Delta\chi^2 = 2.69$ for $\Delta k = 4$ (expected 4 if the extra parameters carry no information, $p = 0.61$); ΔAIC $= -5.31$, ΔBIC $= -19.0$, both favouring SSEE.							
KiDS-1000 published	—	—	260.31	—	—	—	0.759 ± 0.024

The ΛCDM row is the methodological control (R4): the same pipeline, the same raw data, the same nuisance treatment, with the four background parameters $\omega_b, \omega_c, h, n_s$ additionally free. Its purpose is not to provide a new measurement — KiDS-1000 has already published theirs — but to demonstrate that this pipeline reproduces the published ΛCDM result, which is what licenses the SSEE row above it. It does: $S_8 = 0.7571 \pm 0.0194$ against the published 0.759 ± 0.024. We stress that this agreement is *not* a physical measurement of consistency. The published KiDS-1000 value is itself a ΛCDM number —

it is what one obtains by fitting ΛCDM to these data — so a ΛCDM run reproducing it tests the code path, not the cosmology. Reading it as a "tension" would be a category error.

What the control licenses is the comparison in Table 2 between the two rows above the rule, which is like for like: identical data, identical code, identical nuisance treatment. The comparison that carries the weight is the χ^2 block, Eq. (4): SSEE pays 2.69 in χ^2 for giving up four parameters, where four uninformative parameters would have been worth 4. The $\Delta S_8 = 0.0012$ in the right-hand block is a restatement of the same fact in derived quantities — under a tenth of either error bar, and, since both fits are to the *same* data, with strongly correlated uncertainties that make the true discrepancy smaller still. SSEE therefore describes the raw shear data as well as ΛCDM does, using four fewer free parameters and with its background fixed algebraically rather than fitted.

Two entries in the ΛCDM row must not be quoted in isolation: $\sigma_8 = 0.875 \pm 0.089$ and $\Omega_m = 0.231 \pm 0.046$ (and correspondingly the broad $\log A$) are individually poorly determined, because cosmic shear constrains essentially only their combination S_8 — which is why that combination is the reported statistic. The published KiDS-1000 analysis shows the same degeneracy.

5.1 Historical note: the particle's exclusion, for the record

Before the subtraction defining $\Omega_{\phi\mathrm{DM}}$ was recognized as physically empty (§2), the two-sector hypothesis was tested directly against raw KiDS-1000 data with the now-retracted particle mass $m_\phi = 40.70$ eV. Free-streaming suppresses σ_8 by 10.4% at that mass, against a raw-data uncertainty of 3.16%; the amplitude A_s cannot compensate because free-streaming changes the *shape* of $P(k)$, not its overall scale. A measured suppression law,$\Delta\sigma_8(\%) = 52088 \cdot m_\phi^{-2.283}$, places a lower bound $m_\phi > 70.3$ eV for consistency with the shear data — 1.73× the canonical mass, and outside the maximal mass reachable by the model's own algebraic lineage grammar (§2, maximum multiplier 594.28). We record this bound as part of how the retraction was reached (§2), not as a live result: the particle it would constrain does not exist in the current model.

6 $f\sigma_8$: growth rate against raw BOSS DR12 clustering

The clustering counterpart of §4 asks the same question of the other growth probe: with the background held fixed — algebraically in SSEE, at Planck 2018 in ΛCDM — what amplitude does the raw galaxy clustering prefer, and at what χ^2? The test is the χ^2 over the 222 multipole points, *not* $f\sigma_8$: $f\sigma_8$ is an output of the fit, and therefore a product of the model, the same lesson S_8 taught in §2.

Both models carry identical freedom here: one shared amplitude $\ln(10^{10}A_s)$ plus $(b_1, b_2, b_s, \alpha_0, \alpha_2, \mathrm{SN}_0)$ per (z-bin, hemisphere), 37 in total, of which the 18 that enter the model linearly are marginalised in closed form (§3.4) and 19 are sampled. So $\Delta k = 0$: this comparison carries no parsimony credit in either direction. SSEE's advantage in

parameter count lives at the CMB ($k = 2$ against 6) and in the shear likelihood (9 against 13), not here. What this run answers is narrower and cleaner: *with the same nuisance freedom, does the algebraic background fit raw clustering as well as Planck's?*

Table 3: $f\sigma_8$ against raw BOSS DR12 multipoles (R1/R2), LPT to $k_{\max} = 0.20\,h\,\mathrm{Mpc}^{-1}$. Both backgrounds fixed; 222 data points, 37 free parameters, 185 degrees of freedom. Published values are the BOSS DR12 consensus of Alam et al. [2017].

	$z = 0.38$	$z = 0.51$	$z = 0.61$
SSEE	0.4100 ± 0.0206	0.4114 ± 0.0207	0.4091 ± 0.0206
ΛCDM	0.4208 ± 0.0199	0.4192 ± 0.0198	0.4145 ± 0.0196
Alam et al. [2017]	0.497 ± 0.045	0.458 ± 0.038	0.436 ± 0.034
SSEE vs. published	1.76σ	1.08σ	0.68σ
ΛCDM vs. published	1.55σ	0.91σ	0.55σ

The fit quality is the headline:

	$\ln(10^{10}A_s)$	χ^2
ΛCDM	2.7902 ± 0.0945	207.58
SSEE	2.7614 ± 0.1010	205.55

$\Delta\chi^2_{\mathrm{SSEE-\Lambda CDM}} = -2.03$ over 222 points with the same 37 free parameters, i.e. a mild preference for the algebraic background, not a decisive one. Each model carries its own neutrino mass ($\sum m_\nu = 0.06849\,\mathrm{eV}$ in SSEE, 0.06 eV in ΛCDM); an earlier version of these chains ran SSEE with the ΛCDM value and found $\Delta\chi^2 = +0.06$, so the correction moved the comparison towards SSEE, not away from it. The algebraic background describes raw galaxy clustering at least as well as Planck's fitted one.

Method control. Our own ΛCDM row is the control: if the pipeline is sound, it must reproduce what BOSS published. It does, within 1.55σ, 0.91σ and 0.55σ at the three redshifts. We flag two caveats rather than let the reader infer agreement we have not earned. First, $z = 0.38$ at 1.55σ is the weakest of the three and we do not claim it as clean. Second, our uncertainties (± 0.0196 to ± 0.0207) are less than half the published ones (± 0.034 to ± 0.045), because we hold the background *fixed* whereas Alam et al. [2017] marginalise over it: the smaller error bar buys its precision by not asking a question, and the two are therefore not the same quantity. A ΛCDM run with the background free — the version that would make the error bars comparable — is the outstanding item for this section.

Clustering does not measure the amplitude. The marginal means above, 2.7614 ± 0.1010 (SSEE) and 2.7902 ± 0.0945 (ΛCDM), are not a measurement of A_s, and we do not compare them with the CMB. Three measurements show it. First, the marginal and conditional widths of $\log(10^{10}A_s)$ differ by a factor $D = \sigma_{\mathrm{marg}}/\sigma_{\mathrm{cond}} = 4.13$, the amplitude

being degenerate with the linear bias ($r = -0.90$); for KiDS on the same fixed background $D = 1.16$. Second, the central value moves with the estimator: the profile minimum of the same data, same $k_{\max}$ and same model sits at 2.9448 ± 0.1239, 1.82σ from the marginal mean — a volume effect of marginalising over the nuisance parameters, which biases the amplitude low. Third, fixing the amplitude to the CMB value costs the clustering χ^2 only 0.63 ($197.44 \to 198.07$ in the profile). Galaxy clustering is therefore consistent with the CMB amplitude and does not vote on it.

The displacement is not ω_c in disguise. Amplitude and cold dark matter density are degenerate in clustering: both raise the small-scale power. A referee is entitled to ask whether the low $\ln(10^{10}A_s)$ above is really a low ω_c wearing another name. We answer with the symmetric experiment. Holding the amplitude *fixed* at the CMB value and profiling ω_c instead — everything else in the background still fixed, only the galaxy biases fitted — BOSS asks for $\omega_c = 0.109937 \pm 0.003994$, which is 2.40σ from the algebraic $\mathrm{KAL}_0\,\omega_b\,n_s = 0.119514$. Returning the amplitude BOSS itself prefers, ω_c comes back to 0.117450 ± 0.004079, i.e. 0.51σ from the algebraic value. Each point of these profiles minimises the galaxy biases from five starting points and is swept along the grid in both directions; ten independent random starts at the minimum and at both ends of the grid fail to lower χ^2 by more than 0.01. The pre-registered control is that second profile: if ω_c had not returned, the first would have been measuring the edge of the parametrisation rather than the data. It returns.

The same profile run on the ΛCDM background, with identical rigidity, displaces in the same direction and by a comparable amount ($0.114346 \to 0.107801$, i.e. 3.10σ from Planck's $\omega_c = 0.1200$ once the CMB amplitude is imposed), so the displacement is a property of the data and not manufactured by either background. Two consequences follow. First, the low clustering amplitude is not a disguised matter density. Second — and this was not sought — the algebraic ω_c sits *closer* to what a structure survey at $z \simeq 0.5$ asks for (0.51σ) than Planck's fitted ΛCDM value does (1.40σ).

We deliberately do not quote a χ^2 comparison between models with ω_c free. In SSEE that parameter is fixed by algebra: releasing it takes the model outside itself, so such a number would not describe SSEE and cannot arbitrate between the two frameworks. The profile is a diagnostic of where the data pulls, not a fit. What it does license is the statement that ΛCDM gains 4.878 in χ^2 from that (equally illegitimate) freedom against 1.177 for SSEE — Planck's ω_c needs the larger correction. The legitimate model comparison remains the one in the table above, each framework run as itself.

6.1 Why the exploratory linear-Kaiser stage does not publish

An earlier, exploratory pass over the same data used linear Kaiser redshift-space distortions rather than the one-loop model reported above.consensus values (Alam et al. 2017 [Alam et al., 2017]) with the growth rate left free, as in the official analysis, the pipeline reproduces the published values within 0.92σ on average ($\chi^2/\mathrm{dof} = 1.048$; §3.4), confirming

the machinery — data reading, window convolution, and Alcock–Paczynski deformation — is sound. However, the linear model under-predicts $f\sigma_8$ by 9–15% (increasing with z) once matched against small-scale modes ($k > 0.08\,h/\mathrm{Mpc}$), a known limitation of first-order perturbation theory in the mildly non-linear regime. The production run for this paper uses a one-loop effective-theory model (velocileptors, Chen, Vlah & White 2020 [Chen et al., 2020]) that resolves this limitation, and §6 above reports its output, not the linear-theory exploratory result.

7 Conclusions

The growth sector of SSEE has a single matter density, $\Omega_m = 0.308881$, entering both the background and the growth of structure with no partition, no additional particle, and no non-minimal coupling parameter. Measured against raw survey data — rather than against ΛCDM-conditioned summary statistics — it has no S_8 tension. On KiDS-Legacy the amplitude the shear asks for lies 0.49σ from the one the model's own CMB fit supplies; with that amplitude fixed the growth sector has no free cosmological parameter, the fit costs nothing ($\chi^2_{\mathrm{min}} = 417.97$ against 418.34), and $S_8 = 0.8273$ is predicted against 0.8265 ± 0.0176 measured. On the earlier KiDS-1000 release the algebraic background matches ΛCDM run through the same pipeline with four more free parameters ($\Delta\chi^2 = 2.69$, $\Delta\mathrm{BIC} = -19.0$). Against raw BOSS DR12 clustering the algebraic background fits marginally better ($\Delta\chi^2 = -2.03$ at equal freedom, not a decisive difference), and clustering is consistent with the CMB amplitude without constraining it. The primordial amplitude is therefore counted once, at the CMB, and carried unchanged to every late-time probe. Cosmic shear and clustering do not discriminate between the backgrounds; what separates them is the parameter count.

The methodological point (§3) is offered as a general caution independent of SSEE: any model without background-level free parameters cannot exploit the mechanism by which combined datasets average away their own internal tensions, and is therefore a comparatively clean probe of whether those tensions are real.

References

S. Alam et al. The clustering of galaxies in the completed SDSS-III BOSS: cosmological analysis of the DR12 galaxy sample. *Mon. Not. Roy. Astron. Soc.*, 470:2617, 2017. doi: 10.1093/mnras/stx721.

C. Alcock and B. Paczynski. An evolution free test for non-zero cosmological constant. *Nature*, 281:358, 1979. doi: 10.1038/281358a0.

A. Amon et al. Dark energy survey year 3 results: cosmology from cosmic shear and robustness to modeling uncertainty. *Phys. Rev. D*, 105:023514, 2022. doi: 10.1103/PhysRevD.105.023514.

M. Asgari et al. KiDS-1000 cosmology: cosmic shear constraints and comparisons between two point statistics. *Astron. Astrophys.*, 645:A104, 2021. doi: 10.1051/0004-6361/202039070.

D. Baumann, A. Nicolis, L. Senatore, and M. Zaldarriaga. Cosmological non-linearities as an effective fluid. *J. Cosmol. Astropart. Phys.*, 2012(07):051, 2012. doi: 10.1088/1475-7516/2012/07/051.

F. Beutler et al. The clustering of galaxies in the completed SDSS-III BOSS sample: anisotropic galaxy clustering in Fourier-space. *Mon. Not. Roy. Astron. Soc.*, 466:2242, 2017. doi: 10.1093/mnras/stw3298.

S. Bridle and L. King. Dark energy constraints from cosmic shear power spectra: impact of intrinsic alignments on photometric redshift requirements. *New J. Phys.*, 9:444, 2007. doi: 10.1088/1367-2630/9/12/444.

S.-F. Chen, Z. Vlah, and M. White. A new analysis of galaxy 2-point functions in the BOSS survey, including full-shape information and post-reconstruction BAO. *J. Cosmol. Astropart. Phys.*, 2020(02):008, 2020. doi: 10.1088/1475-7516/2020/02/008.

M. Chevallier and D. Polarski. Accelerating universes with scaling dark matter. *Int. J. Mod. Phys. D*, 10:213, 2001. doi: 10.1142/S0218271801000822.

E. Di Valentino et al. Cosmology intertwined III: $f\sigma_8$ and s_8. *Astropart. Phys.*, 131:102604, 2021. doi: 10.1016/j.astropartphys.2021.102604.

F. Feroz, M. P. Hobson, and M. Bridges. MultiNest: an efficient and robust Bayesian inference tool for cosmology and particle physics. *Mon. Not. Roy. Astron. Soc.*, 398:1601, 2009. doi: 10.1111/j.1365-2966.2009.14548.x.

C. Heymans et al. KiDS-1000 cosmology: multi-probe weak gravitational lensing and spectroscopic galaxy clustering constraints. *Astron. Astrophys.*, 646:A140, 2021. doi: 10.1051/0004-6361/202039063.

E. V. Linder. Exploring the expansion history of the universe. *Phys. Rev. Lett.*, 90:091301, 2003. doi: 10.1103/PhysRevLett.90.091301.

A. J. Mead, J. A. Peacock, C. Heymans, S. Joudaki, and A. F. Heavens. An accurate halo model for fitting non-linear cosmological power spectra and baryonic feedback models. *Mon. Not. Roy. Astron. Soc.*, 454:1958, 2015. doi: 10.1093/mnras/stv2036.

Planck Collaboration. Planck 2018 results. VI. cosmological parameters. *Astron. Astrophys.*, 641:A6, 2020. doi: 10.1051/0004-6361/201833910.

J. Torrado and A. Lewis. Cobaya: Code for Bayesian Analysis of hierarchical physical models, 2021.

M. J. Wilson, J. A. Peacock, A. N. Taylor, and S. de la Torre. Rapid modelling of the redshift-space power spectrum multipoles for a masked density field. *Mon. Not. Roy. Astron. Soc.*, 464:3121, 2017. doi: 10.1093/mnras/stw2576.

A. H. Wright et al. KiDS-Legacy: Cosmological constraints from cosmic shear with the complete Kilo-Degree Survey. *Astron. Astrophys.*, 703:A158, 2025.